

# Intel® Cloud Optimization Modules for Microsoft Azure\*: Stable Diffusion Distributed Training

This module illustrates the process of fine-tuning a stable diffusion model on 3rd or 4th Generation [Intel® Xeon® Scalable Processors](#) on [Microsoft Azure](#) in a distributed architecture. We will employ the [Stable Diffusion v1](#) implementation and use the [Dicoo images](#) from Hugging Face\* Hub.

To capitalize on the capabilities of Intel hardware, we also integrate [Intel AI software](#) offerings, like the [Intel® Extension for PyTorch\\*](#) and the [Intel® oneAPI Collective Communications Library \(oneCCL\)](#).

## Azure Cluster Setup

For distributed fine-tuning, our example deploys a cluster consisting of three 3<sup>rd</sup> Gen Intel® Xeon® CPUs from the [Azure Dv5 series](#) with an Ubuntu 22.04 image and 32 GiB of storage deployed on an [Azure Trusted Virtual Machine](#). For maximum performance, we recommend using the precision data type **bfloat16** on 4<sup>th</sup> Gen Intel® Xeon® CPUs on Azure with the deep learning acceleration engine, [Intel® Advanced Matrix Extensions \(Intel® AMX\)](#).

If you are using a 4<sup>th</sup> Gen Xeon CPU, you can verify AMX support by running:

```
lscpu | grep amx
```

You should see the following flags:

```
amx_bf16 amx_tile amx_int8
```

## Intel® Extension for PyTorch\*

Intel upstreams as many optimizations as possible to PyTorch. These features, however, often debut in the Intel® Extension for PyTorch. Install PyTorch ([guide](#)), and then the Intel Extension for PyTorch:

```
pip install intel_extension_for_pytorch==2.0.0
```

To enable the optimizations, only add these **two lines** to your Python\* code:

```
import intel_extension_for_pytorch as ipex

model.train()
model, optimizer = ipex.optimize(model, optimizer=optimizer,
                                dtype="torch.bfloat16", level="O1", inplace=True)
```

[Documentation](#)[Cheat Sheet](#)[Examples](#)[Tuning Guide](#)

## Install oneCCL

Download the appropriate wheel file and install it using the following commands:

```
wget https://intel-extension-for-pytorch.s3.amazonaws.com/torch_ccl/cpu/oneccl_bind_pt-2.0.0%2Bcpu-cp38-cp38-linux_x86_64.whl
pip install oneccl_bind_pt-2.0.0+cpu-cp38-cp38-linux_x86_64.whl
```

To use `oneccl_bindings_for_pytorch`, source the environment by running the following command:

```
oneccl_path=$(python -c "from oneccl_bindings_for_pytorch import cwd; print(cwd)")
source $oneccl_path/env/setvars.sh
```

To launch distributed fine-tuning:

```
mpirun -f ~/hosts -n 3 -ppn 1 accelerate launch textual_inversion.py --
pretrained_model_name_or_path="runwayml/stable-diffusion-v1-5" --train_data_dir="dicoo" ...
```

The `mpirun` command runs the fine-tuning process across 3 machines (`-n 3`) with one process per machine (`-ppn 1`). Additionally, environment variables can be set using `-genv` argument.

Next Steps:

[All Cloud Modules](#) | [GitHub Repo](#) | [DevMesh Discord](#)

\*Names and brands may be claimed as the property of others.